

Capstone_Project [1]: MovieLens report

Abedulla M. A. El-Saidy

12/12/2021

1. Introduction

Data Science: Capstone, is the final course in the HarvardX Professional Certificate in Data Science. For this project, I will create a movie recommendation system using the MovieLens dataset.

I will create my own recommendation system using all the tools that have been shown throughout the courses in this series. I will use the 10M version of the MovieLens dataset to make the computation a little easier. I will download the MovieLens data and run code provided to generate my datasets. I will train a machine learning algorithm using the inputs to predict movie ratings.

2. Loading Data

#Loading Data from MovieLens Dataset zip file,

Note: this process could take a couple of minutes

```
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
```

```
library(tidyverse) library(caret) library(data.table)
```

MovieLens 10M dataset:

<https://grouplens.org/datasets/movielens/10m/>

<http://files.grouplens.org/datasets/movielens/ml-10m.zip>

```
dl <- tempfile() download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
```

```
ratings <- fread(text = gsub(":", ",", readLines(unzip(dl,"ml-10M100K/ratings.dat"))), col.names =
c("userId", "movieId", "rating", "timestamp"))
```

```
movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\t", 3)
colnames(movies) <- c("movieId", "title", "genres")
```

```

if using R 3.6 or earlier:

movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId], title =
as.character(title), genres = as.character(genres)) # if using R 4.0 or later: movies <-
as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId), title = as.character(title),
genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

Validation set will be 10% of Movielens data

set.seed(1, sample.kind="Rounding") # if using R 3.5 or earlier, use set.seed(1) test_index <-
createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE) edx <- movielens[-
test_index,] temp <- movielens[test_index,]

Make sure userId and movieId in validation set are also in edx set

validation <- temp %>% semi_join(edx, by = "movieId") %>% semi_join(edx, by = "userId")

Add rows removed from validation set back into edx set

removed <- anti_join(temp, validation) edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

#then I store it in file called "gresdy.RData"

```

3. Data exploration

```

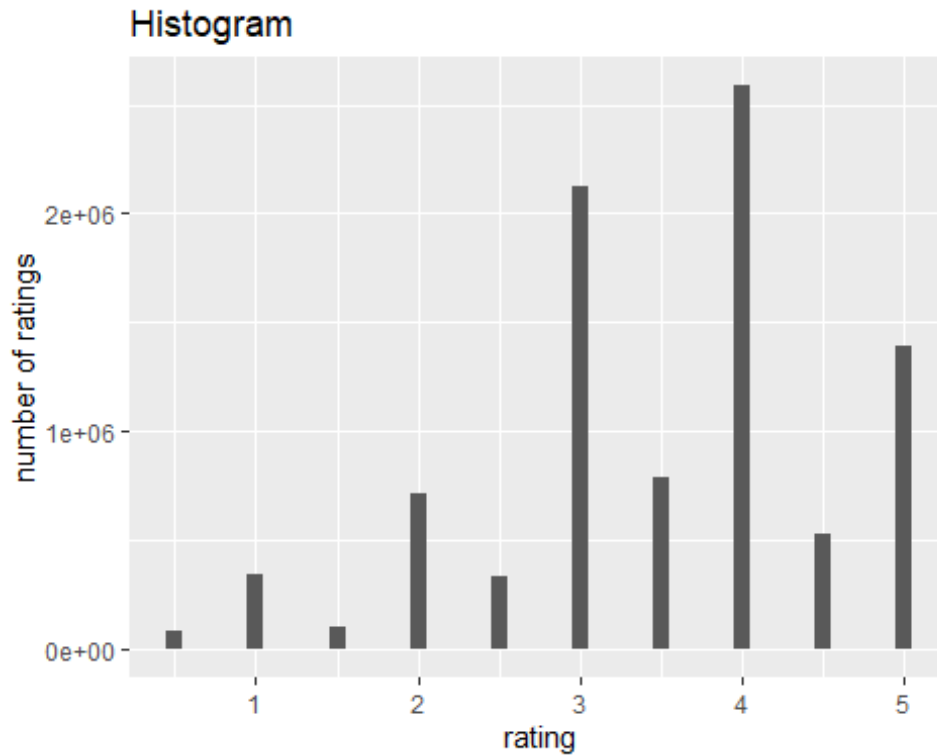
#Data exploration

summary(edx)

##      userId      movieId      rating      timestamp
## Min.   :    1   Min.   :    1   Min.   :0.500   Min.   :7.897e+08
## 1st Qu.:18124   1st Qu.:   648   1st Qu.:3.000   1st Qu.:9.468e+08
## Median :35738   Median :  1834   Median :4.000   Median :1.035e+09
## Mean   :35870   Mean   :  4122   Mean   :3.512   Mean   :1.033e+09
## 3rd Qu.:53607   3rd Qu.:  3626   3rd Qu.:4.000   3rd Qu.:1.127e+09
## Max.   :71567   Max.   :65133   Max.   :5.000   Max.   :1.231e+09
##      title      genres
## Length:9000055   Length:9000055
## Class :character Class :character
## Mode  :character Mode  :character
##
# histogram of ratings

ggplot(edx, aes(x= edx$rating)) + geom_histogram( binwidth = 0.1) +
labs(x="rating", y="number of ratings") + ggtitle("Histogram")

```



```
# top 10 movies
top_10 <- edx %>% group_by(title) %>% summarize(count=n()) %>%
  top_n(10,count) %>% arrange(desc(count))
print(top_10)
```

```
## # A tibble: 10 x 2
##   title                                coun
##   <chr>                                <int>
## 1 Pulp Fiction (1994)                  3136
## 2 Forrest Gump (1994)                  3107
## 3 Silence of the Lambs, The (1991)    3038
## 4 Jurassic Park (1993)                 2936
## 5 Shawshank Redemption, The (1994)    2801
## 6 Braveheart (1995)                   2621
## 7 Fugitive, The (1993)                 2599
## 8 Terminator 2: Judgment Day (1991)    2598
## 9 Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977) 2567
```


2

10 Apollo 13 (1995)

2428

4

```
ggplot(top_10, aes(x=reorder(title,count),y=count))+geom_bar(stat='identity', fill="cyan", width=0.2)+ coord_flip()+ylab("Count")+xlab("Title")
```



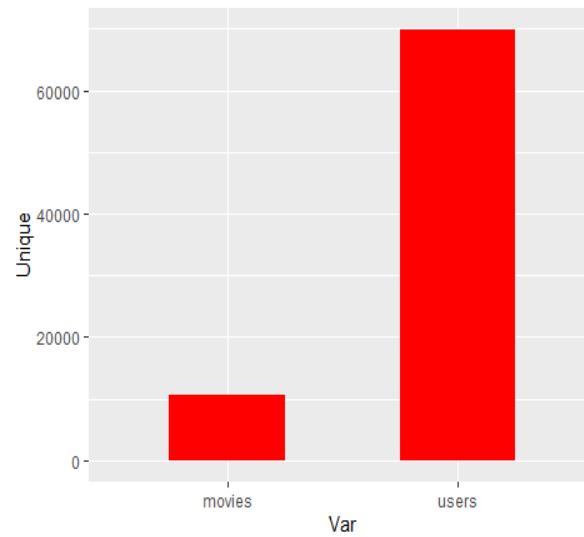
#The number of unique values for movieId and userId

```
unique=data.frame(movies=length(unique(edx$movieId)),users=length(unique(edx$userId)))  
unique
```


```
##   movies users
```

```
## 1  10677 69878
```


```
tunique=as.data.frame( t(unique))  
ggplot(tunique, aes(x=row.names(tunique),y=tunique[,1]))+geom_bar(stat='identity', fill="red", width=.5)+xlab("Var")+ylab("Unique")
```



```
# histogram movie Versus ratings
edx %>% count(movieId) %>% ggplot(aes(n)) + geom_histogram(bins=20,color="Green")+ scale_x_log10() +xlab("Mov ID")+ylab("# of rates")
```



```
# histogram User Versus ratings
edx %>% count(userId) %>% ggplot(aes(n)) + geom_histogram(bins=20,color="Green")+ scale_x_log10() +xlab("User ID")+ylab("# of rates")
```




#some movies get rated more than others, and some users are more active than others

the relation between rating versus year

```
time=round_date(as_datetime(edx$timestamp), unit = "year")
timedat=data.frame(time,edx$rating)
timedat %>% group_by(time) %>% summarize(edx.rating = mean(edx.rating))
) %>% ggplot(aes(time, edx.rating)) + geom_smooth( colour="red", size
=0.5)
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'



#When compared to current times, the rating was high in the past.

4. methods/analysis

```
# the RMSE function -library(caret) -> RMSE()

#1.simple methods
#The mean of all movies
meo=mean(edx$rating)
mean_only_RMSE=RMSE(validation$rating,meo)
mean_only_RMSE

## [1] 1.061202

# movie effect
e_mov_mean = edx %>% group_by(movieId) %>% summarize(e_mov_avg = mean
(rating-meo))
pred_rat_eachmov = validation %>% left_join(e_mov_mean, by='movieId') %
>% mutate(pr1 = meo + e_mov_avg)
mean_each_movie_RMSE = RMSE(validation$rating,pred_rat_eachmov$pr1)
mean_each_movie_RMSE


## [1] 0.9439087

# user+movie effect
e_user_mean = edx %>% left_join(e_mov_mean, by = "movieId") %>% group_
by(userId) %>% summarize(e_user_avg = mean(rating-meo-e_mov_avg))
pred_rat_eachusr = validation %>% left_join(e_mov_mean, by='movieId') %
>%left_join(e_user_mean, by='userId') %>% mutate(pr2 = meo + e_mov_avg
+ e_user_avg)
mean_each_user_RMSE = RMSE(validation$rating,pred_rat_eachusr$pr2)
mean_each_user_RMSE

## [1] 0.8653488

#Regularization
Lamda <- seq(1, 10, 1)
RMS-Regularization <- sapply(Lamda, function(Lam){
  meo=mean(edx$rating)
  e_mov_avg <- edx %>% group_by(movieId) %>% summarize(e_mov_avg = sum
(rating - meo)/(n()+Lam))
  e_user_mean <- edx %>% left_join(e_mov_avg, by="movieId") %>% group_b
y(userId) %>% summarize(e_user_mean = sum(rating - e_mov_avg - meo)/(n(
)+Lam))
  pred_rat <- validation %>%left_join(e_mov_avg, by = "movieId") %>%lef
t_join(e_user_mean, by = "userId") %>%mutate(pr = meo + e_mov_avg + e_u
ser_mean)

  return(RMSE(validation$rating, pred_rat$pr))
})
plot(Lamda, RMS-Regularization)
```



```

min(RMS_Regularization)

## [1] 0.8648177

#2.Caret Prediction
#sampling
set.seed(1000, sample.kind="Rounding")

## Warning in set.seed(1000, sample.kind = "Rounding"): non-uniform 'Rounding'
## sampler used

miniedx=edx[sample(9000055,10000),]

#Split the mini edx
indx = createDataPartition(miniedx$rating, times=1, p=0.8,list=FALSE)
train = miniedx[indx,]
test = miniedx[-indx,]

#Machine Learning Models
#Linear Model
lm = train(rating~userId+movieId+timestamp, data=train, method="lm")

## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind =
NULL) :
## non-uniform 'Rounding' sampler used

lm

```



```

## Linear Regression
##
## 8001 samples
##    3 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 8001, 8001, 8001, 8001, 8001, 8001, ...
## Resampling results:
##
##    RMSE      Rsquared    MAE
##    1.055194  0.002703584  0.8524619
##
## Tuning parameter 'intercept' was held constant at a value of TRUE

pred1 = predict(lm, test)
lm_RMSE = RMSE(test$rating, pred1 )
lm_RMSE

## [1] 1.048852

#Decision Tree
TRE = train(rating~userId+movieId+timestamp, data=train, method="rpart"
)

## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind =
NULL) :
## non-uniform 'Rounding' sampler used

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
trainInfo, :
## There were missing values in resampled performance measures.

(TRE)

## CART
##
## 8001 samples
##    3 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 8001, 8001, 8001, 8001, 8001, 8001, ...
## Resampling results across tuning parameters:
##
##    cp          RMSE      Rsquared    MAE
##    0.002365630  1.053217  0.01499636  0.8425199
##    0.002425463  1.052815  0.01501857  0.8425324
##    0.010089754  1.052489  0.01723323  0.8471947
##

```

```

## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was cp = 0.01008975.

pred2 = predict(TRE, test)
TRE_RMSE = RMSE(test$rating, pred2)
TRE_RMSE

## [1] 1.049581

#k-Nearest Neighbors
knn = train(rating~userId+movieId+timestamp, data=train, method="knn")


## Warning in (function (kind = NULL, normal.kind = NULL, sample.kind =
NULL) :
## non-uniform 'Rounding' sampler used

knn

## k-Nearest Neighbors
##
## 8001 samples
##    3 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 8001, 8001, 8001, 8001, 8001, 8001, ...
## Resampling results across tuning parameters:
##
##    k  RMSE      Rsquared    MAE
##    5  1.225627  0.0007518559  0.9704332
##    7  1.186605  0.0006304272  0.9414621
##    9  1.161674  0.0004954929  0.9210902
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 9.

plot(knn)

```



```
pred3 = predict(knn, test)
knn_RMSE = RMSE(test$rating, pred3)
knn_RMSE

## [1] 1.100773
```

5. Results

```
Rmse_Result=data.frame(method=c("mean_only", "mean_each_movie", "mean_eac
h_user", "Regularization", "Linear Model", "Decision Tree", "k-Nearest Neig
hbors" ), RMSE=c(mean_only_RMSE, mean_each_movie_RMSE, mean_each_user_RMS
E, min(RMS-Regularization), lm_RMSE, TRE_RMSE, knn_RMSE))
arrange(Rmse_Result, (RMSE))

##           method      RMSE
## 1 Regularization 0.8648177
## 2 mean_each_user 0.8653488
## 3 mean_each_movie 0.9439087
## 4 Linear Model 1.0488522
## 5 Decision Tree 1.0495815
## 6 mean_only 1.0612018
## 7 k-Nearest Neighbors 1.1007733
```

6. Conclusion

The final RMSE is 0.8648177. I constructed and tested numerous models and found that Regularization provided the best accuracy.

7. References

Irizzary,R.,Introduction to Data Science